

Quadrature-swap savings across EquiformerV2 architectures

Architecture-level forward (random-init), $n = 3 \text{ runs} \times 20 \text{ forwards}$ on distinct QM9 batches each, NVIDIA A100-SXM4-40GB

