

Quadrature-swap savings across EquiformerV2 architectures

Architecture-level forward (random-init), $n = 3$ runs $\times$ 20 forwards on distinct QM9 batches each, NVIDIA A100-SXM4-40GB

